
emata

Don't Get Lost in the Random Forests: A Beginner's Guide

Moses Bomera, Emata

April 29th, 2021

Moses Bomera

Data Scientist
Emata (& Laboremus)

Works on alternative credit scoring for microfinance institution Emata, a spin-off of fintech company Laboremus.

Prior to joining Laboremus, worked on research in computer vision and natural language processing at netLabs!UG Research Centre, Makerere University.

We are on a mission to provide affordable digital loans to millions of farmers in East Africa

Today, agri-financing in Africa does NOT work.

Farmers are risky clients

NO (CREDIT) HISTORY,
NO COLLATERAL

Farmers need small loans

HIGH OPERATING COSTS

50-100%

ANNUAL
INTEREST RATE

What does Emata do?

1. Digitise agricultural cooperatives to get data

2. Turn data into credit scoring

3. Offer digital and affordable loans to farmers


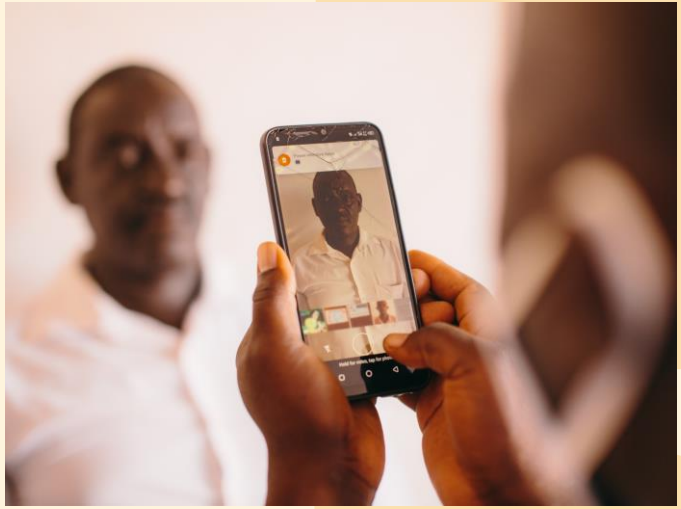


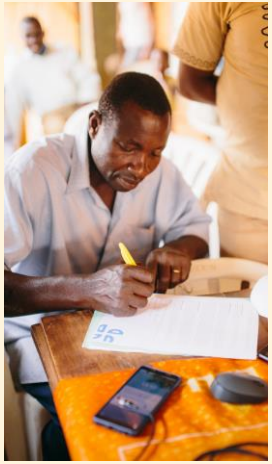
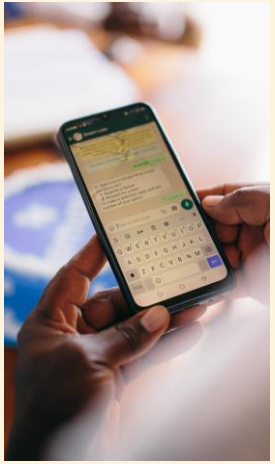
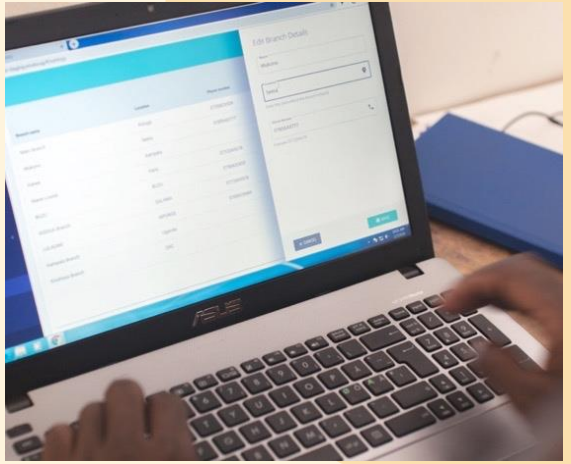


Agenda

- What is a Random Forest?
- Principal component - A Decision Tree
- Calculating the gini value
- The CART Training Algorithm
- Regression
- Gini impurity or Entropy?
- How does it work?
- Regularization Hyperparameters
- TLDR
- Reference Material
- Tutorial

emata

What is a Random Forest?

What is a Random Forest?

It is an ensemble of decision trees.

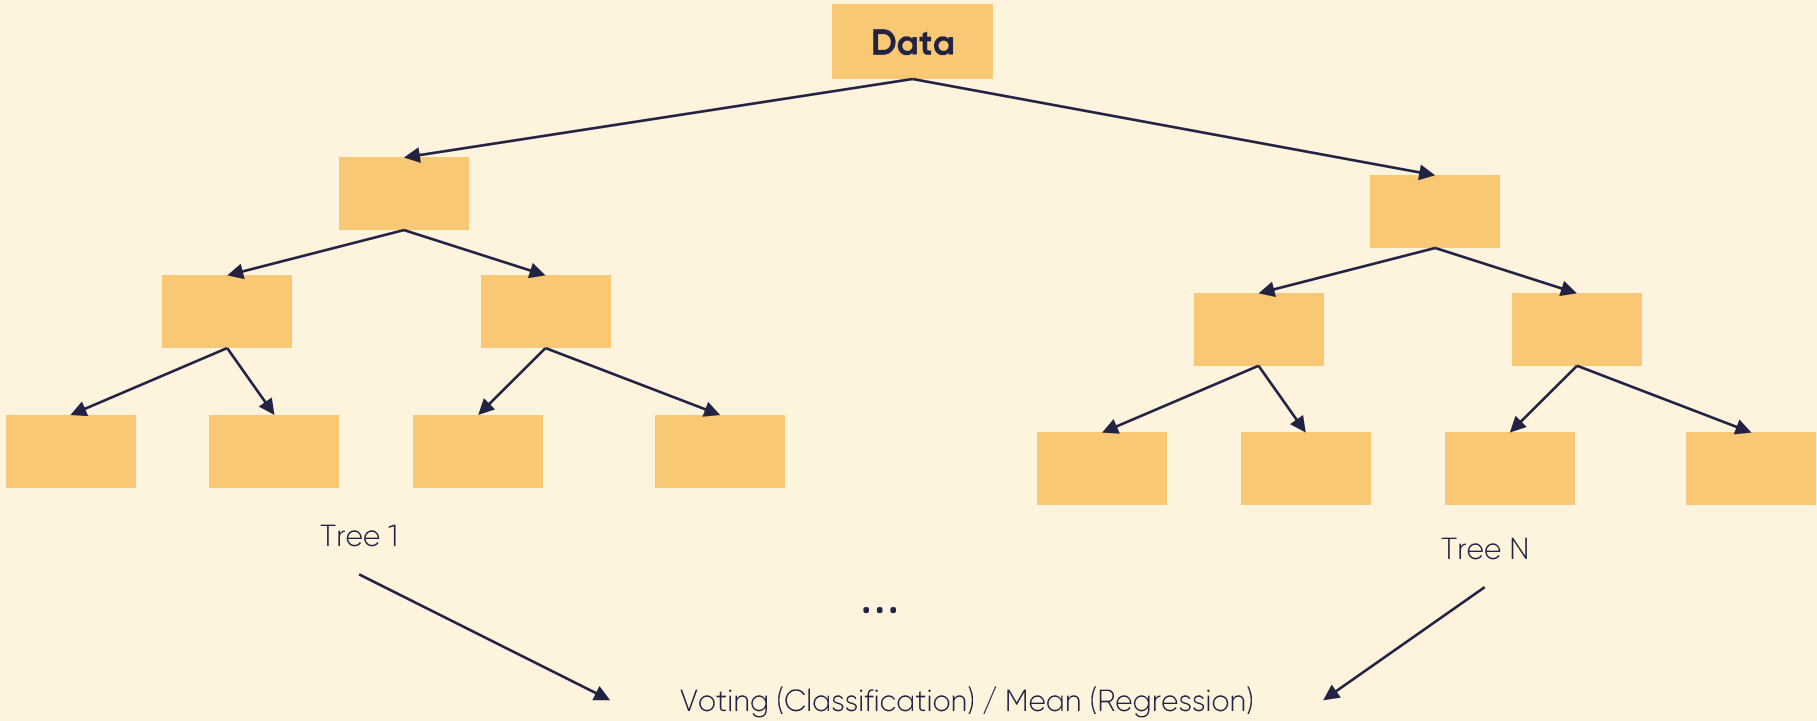

What have Random Forests been used for?

Classification, regression, time series forecasting etc.

Image Classification	Classification + Localization	Object Detection
	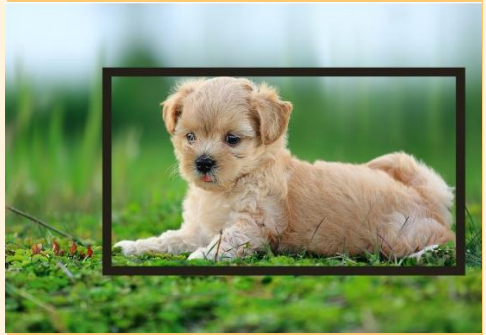	
DOG	DOG	CAT, DOG

Photo source: morioh.com

Random Forests

A decision tree (1/3)

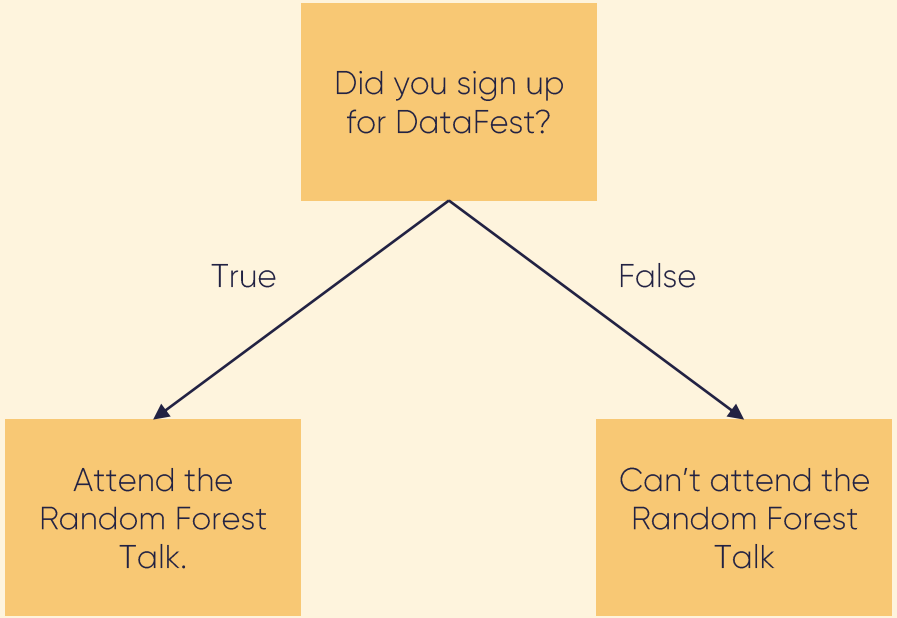

Random Forests

A decision tree (2/3)

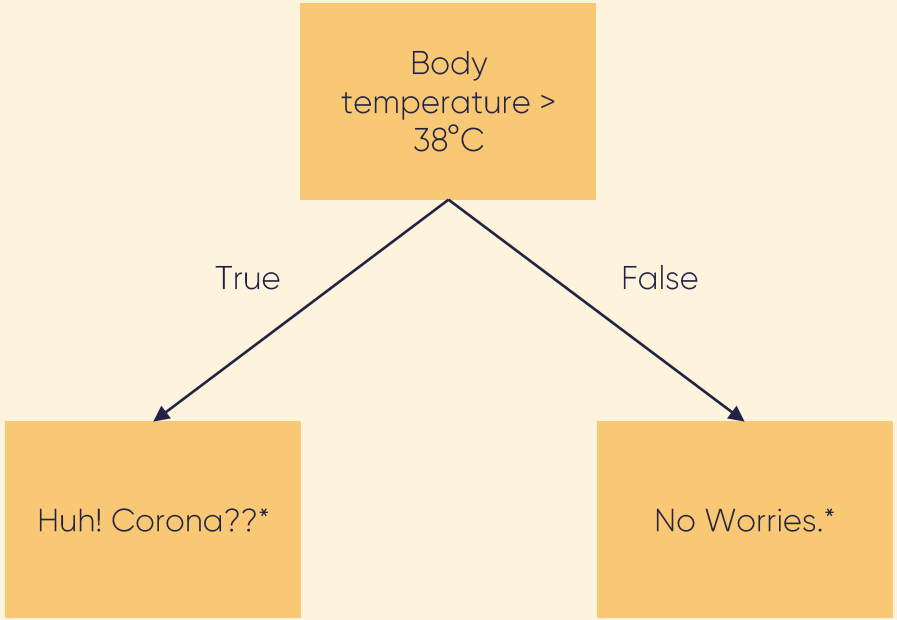

*For tutorial purpose only.

Random Forests

A decision tree (3/3)

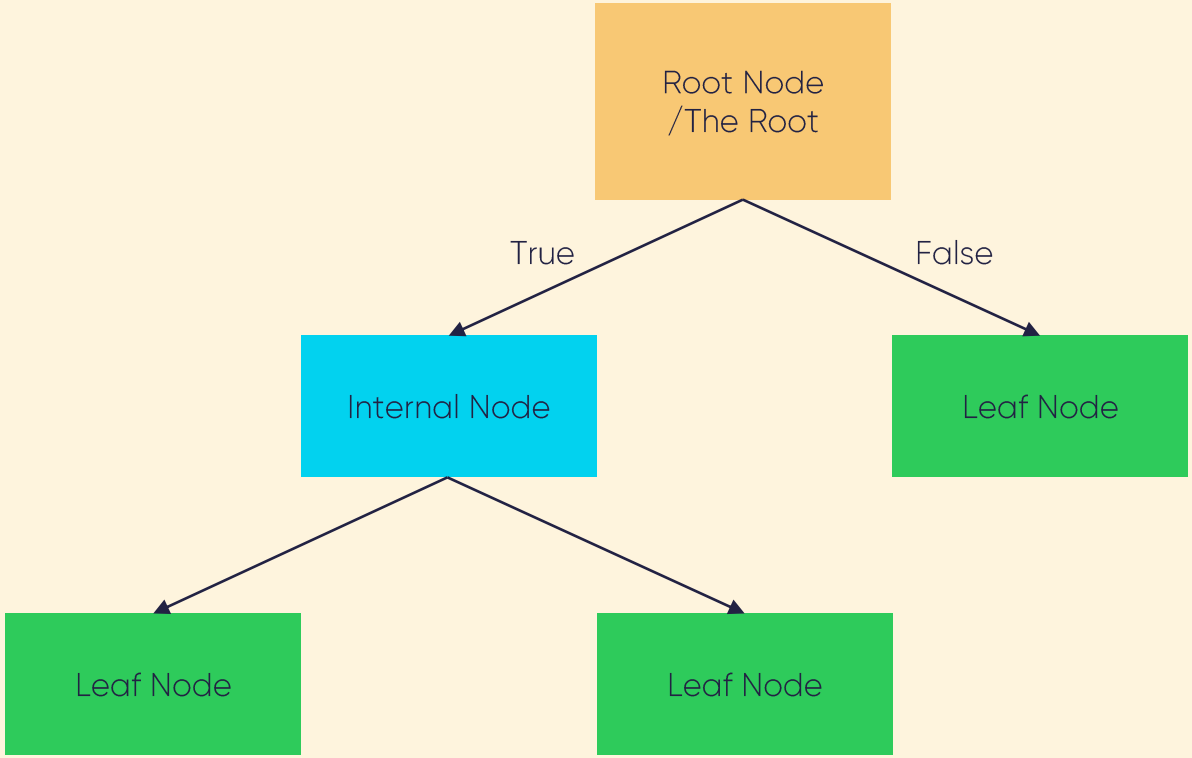

Dataset

petal length	petal width	target
1.4	0.2	setosa
1.3	0.2	setosa
...	...	...
5.1	1.8	virginica

[Iris dataset]

Left; Iris setosa,
Top; Iris virginica,
Bottom; Iris veriscolor

The Theory

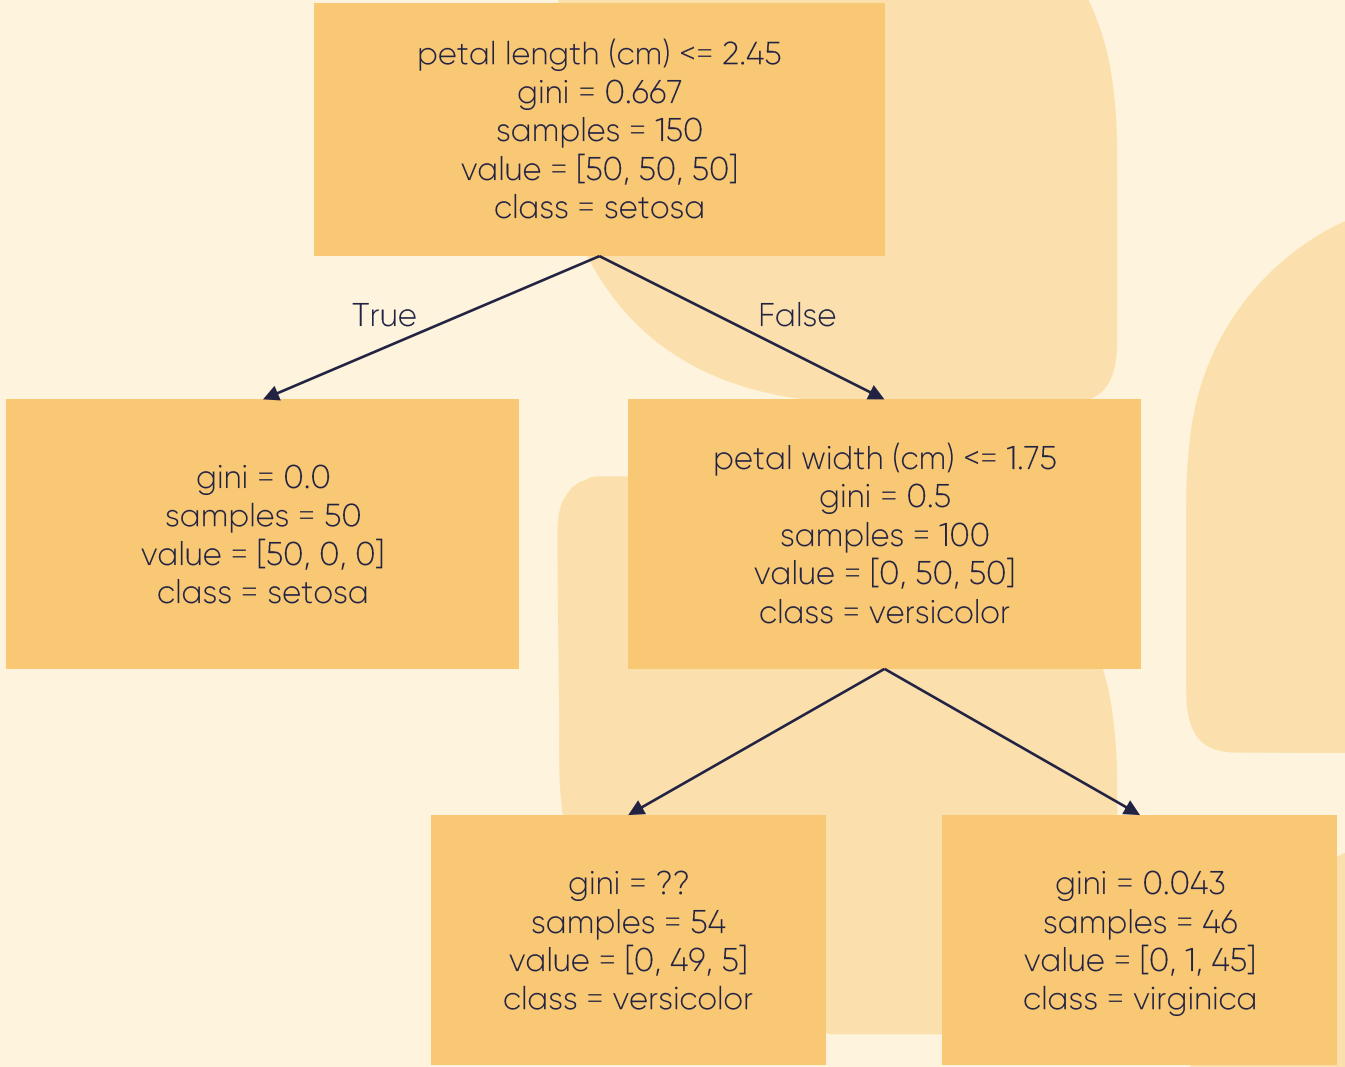
petal length (cm) <= 2.45
gini = 0.667
samples = 150
value = [50, 50, 50]
class = setosa
True
False
gini = 0.0
samples = 50
value = [50, 0, 0]
class = setosa
petal width (cm) <= 1.75
gini = 0.5
samples = 100
value = [0, 50, 50]
class = versicolor
gini = ??
samples = 54
value = [0, 49, 5]
class = versicolor
gini = 0.043
samples = 46
value = [0, 1, 45]
class = virginica

Node Features

Samples: the number of training instances that particular node applies to.

Gini: measures how pure (or impure) a node is.

Class: the target represented by a given node.

$p_{i,k}$ is the ratio of class k instances to all the training instances in i^{th} node.

gini = 0.043
samples = 46
value = [0, 1, 45]
class = virginica

$$G_i = 1 - \sum_{k=1}^{n} p_{i,k}^2$$

Calculating the gini value

gini = ??
samples = 54
value = [0, 49, 5]
class = versicolor

$$G_i = 1 - \sum_{k=1}^{n} {p_{i,k}}^2$$

$$1 - \left(\frac{0}{54}\right)^2 - \left(\frac{49}{54}\right)^2 - \left(\frac{5}{54}\right)^2 =$$

Classification and Regression Tree (CART) Algorithm

$$J(k, t_k) = \frac{m_{\text{left}}}{m} G_{\text{left}} + \frac{m_{\text{right}}}{m} G_{\text{right}}$$

where $\begin{cases} G_{\text{left/right}} \text{ measures the impurity of the left/right subset,} \\ m_{\text{left/right}} \text{ is the number of instances in the left/right subset.} \end{cases}$

1. CART splits the training set into two subsets using a single feature, k and a threshold purity, t_k.
2. The selection process involves minimizing the cost function
3. 1 and 2 are repeated recursively for each subset until the maximum depth is reached.

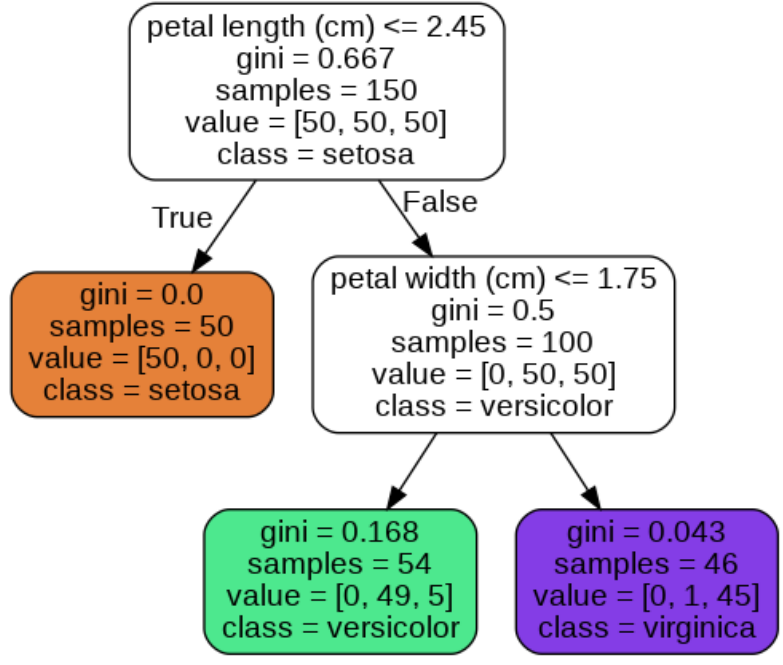

Regression

- Similar to the classification approach.
- The prediction is the average of the samples associated with the leaf node.
- For the gini impurity split, the regression tree uses mean squared error i.e. attempts to minimize the mse.

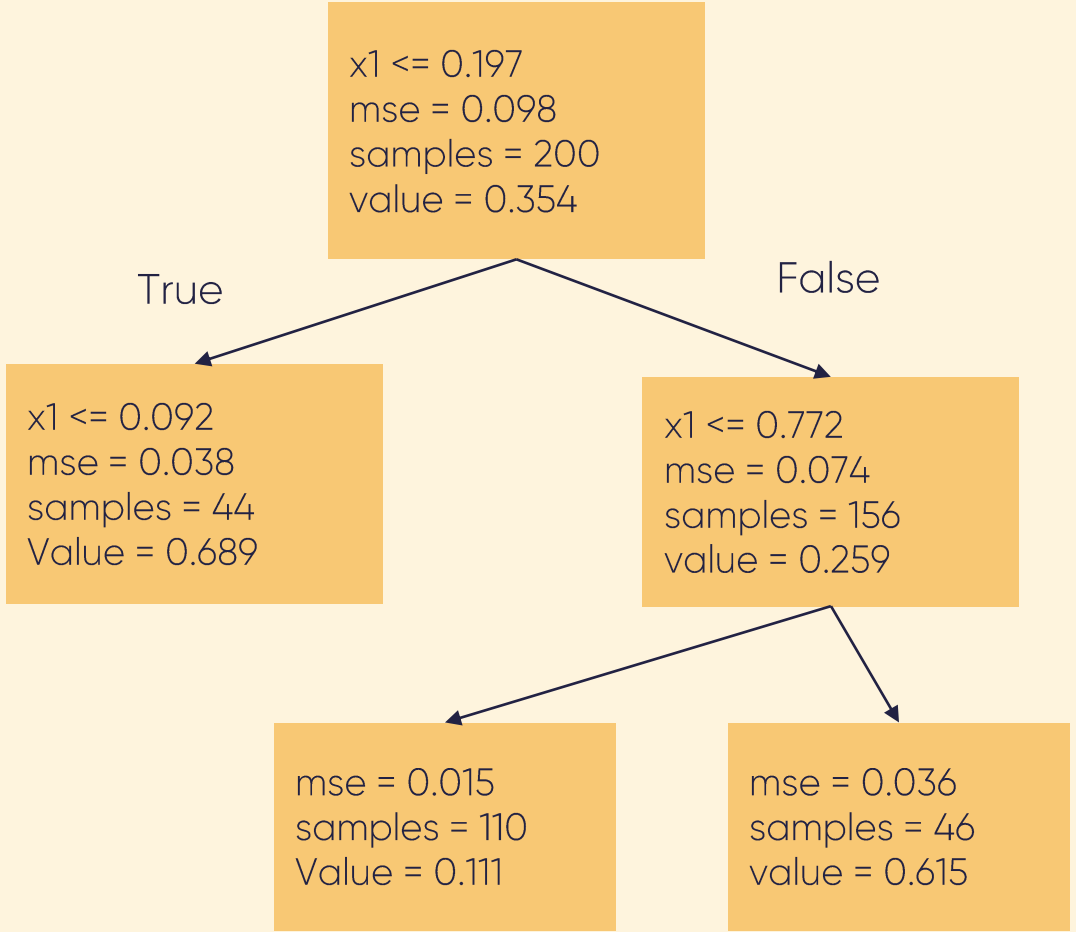

Gini impurity or Entropy?

- **While gini is the go-to impurity measure**, you can also use entropy as a measurement.
- Entropy is based on the same concept from thermodynamics.
- Entropy approaches zero when molecules are still and well ordered.
- In machine learning, entropy is zero when a set contains instances of only one class.

$$H_i = -\sum_{\substack{k=1 \\ p_{i,k} \neq 0}}^{n} p_{i,k} \log_2 \left(p_{i,k}\right)$$

What is a Random Forest?

It is an ensemble of decision trees.

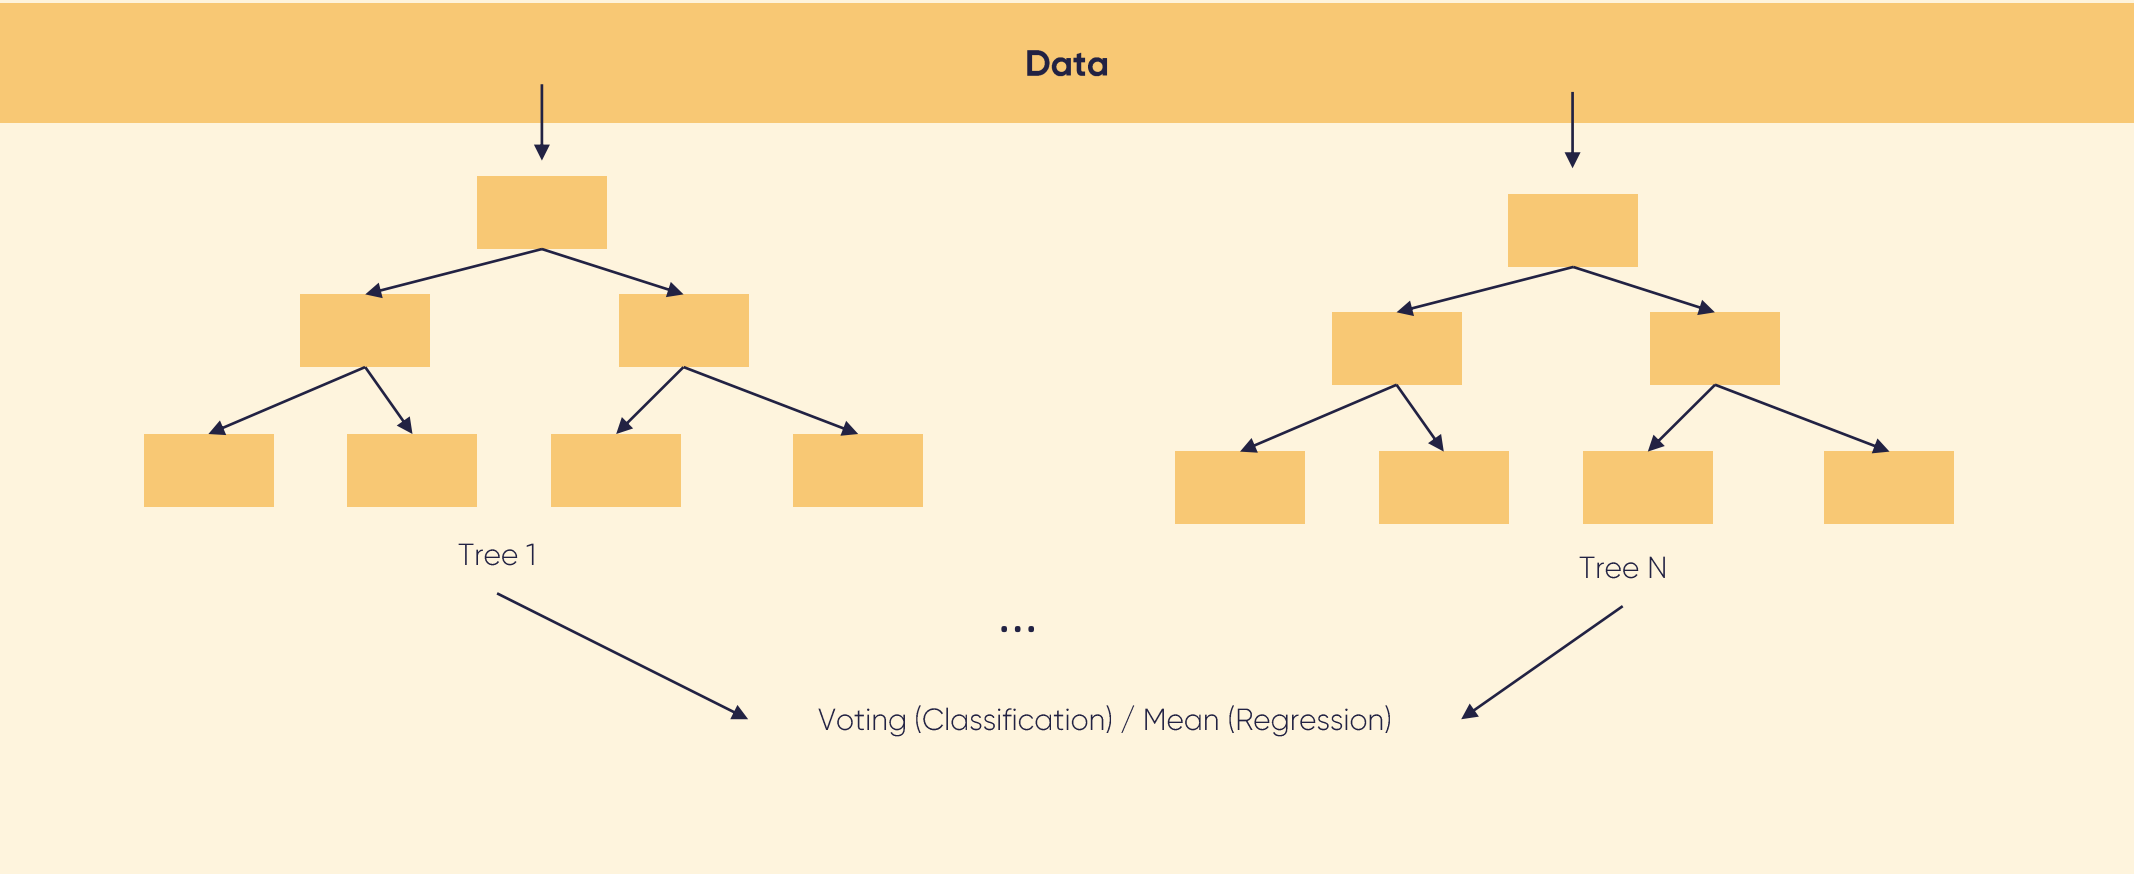

How does it work? (1/2)

- A random forest is an ensemble of decision trees.
- Trained usually through the bagging method (or pasting).
- Each tree in a random forest gives a prediction, for classification, the class with the majority of votes is the prediction.
- In regression, the average of each tree's prediction is the prediction.

How does it work? (2/2)

- RF introduces randomness when growing trees.
- Doesn't make a split using the very best feature, instead it selects the best feature among a random subset of features.
- Results in greater tree diversity.
- Trading a higher bias with lower variance yielding a better model than an individual decision tree.

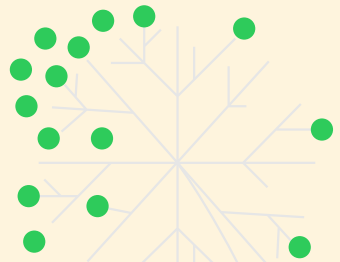

Regularization Hyperparameters

- *n_estimators* - the number of trees to use.
- *max_depth* – the depth of the tree.
- *min_samples_split* - the minimum number of samples an internal node must have before it can be split.
- *min_samples_leaf* - the minimum number of samples required to be at a leaf node
- *max_leaf_nodes* - the maximum number of leaf nodes.
- *max_features* - the maximum number of features that are evaluated for splitting at each node.

Increasing *min_** hyperparameters or reducing *max_** hyperparameters will regularize the model.

TLDR;

PROS of using Random Forests

- Random Forests can be used for both classification and regression.
- Can handle large datasets with high dimensionality, (can be used for dimensionality reduction).
- Works well with unscaled datasets.

CONS of using Random Forests

- It is a high variance model, so it is important to regularize the model when training.
- Given the numerous hyper-parameters, and the number of trees, the decision process becomes a bit of a black box.

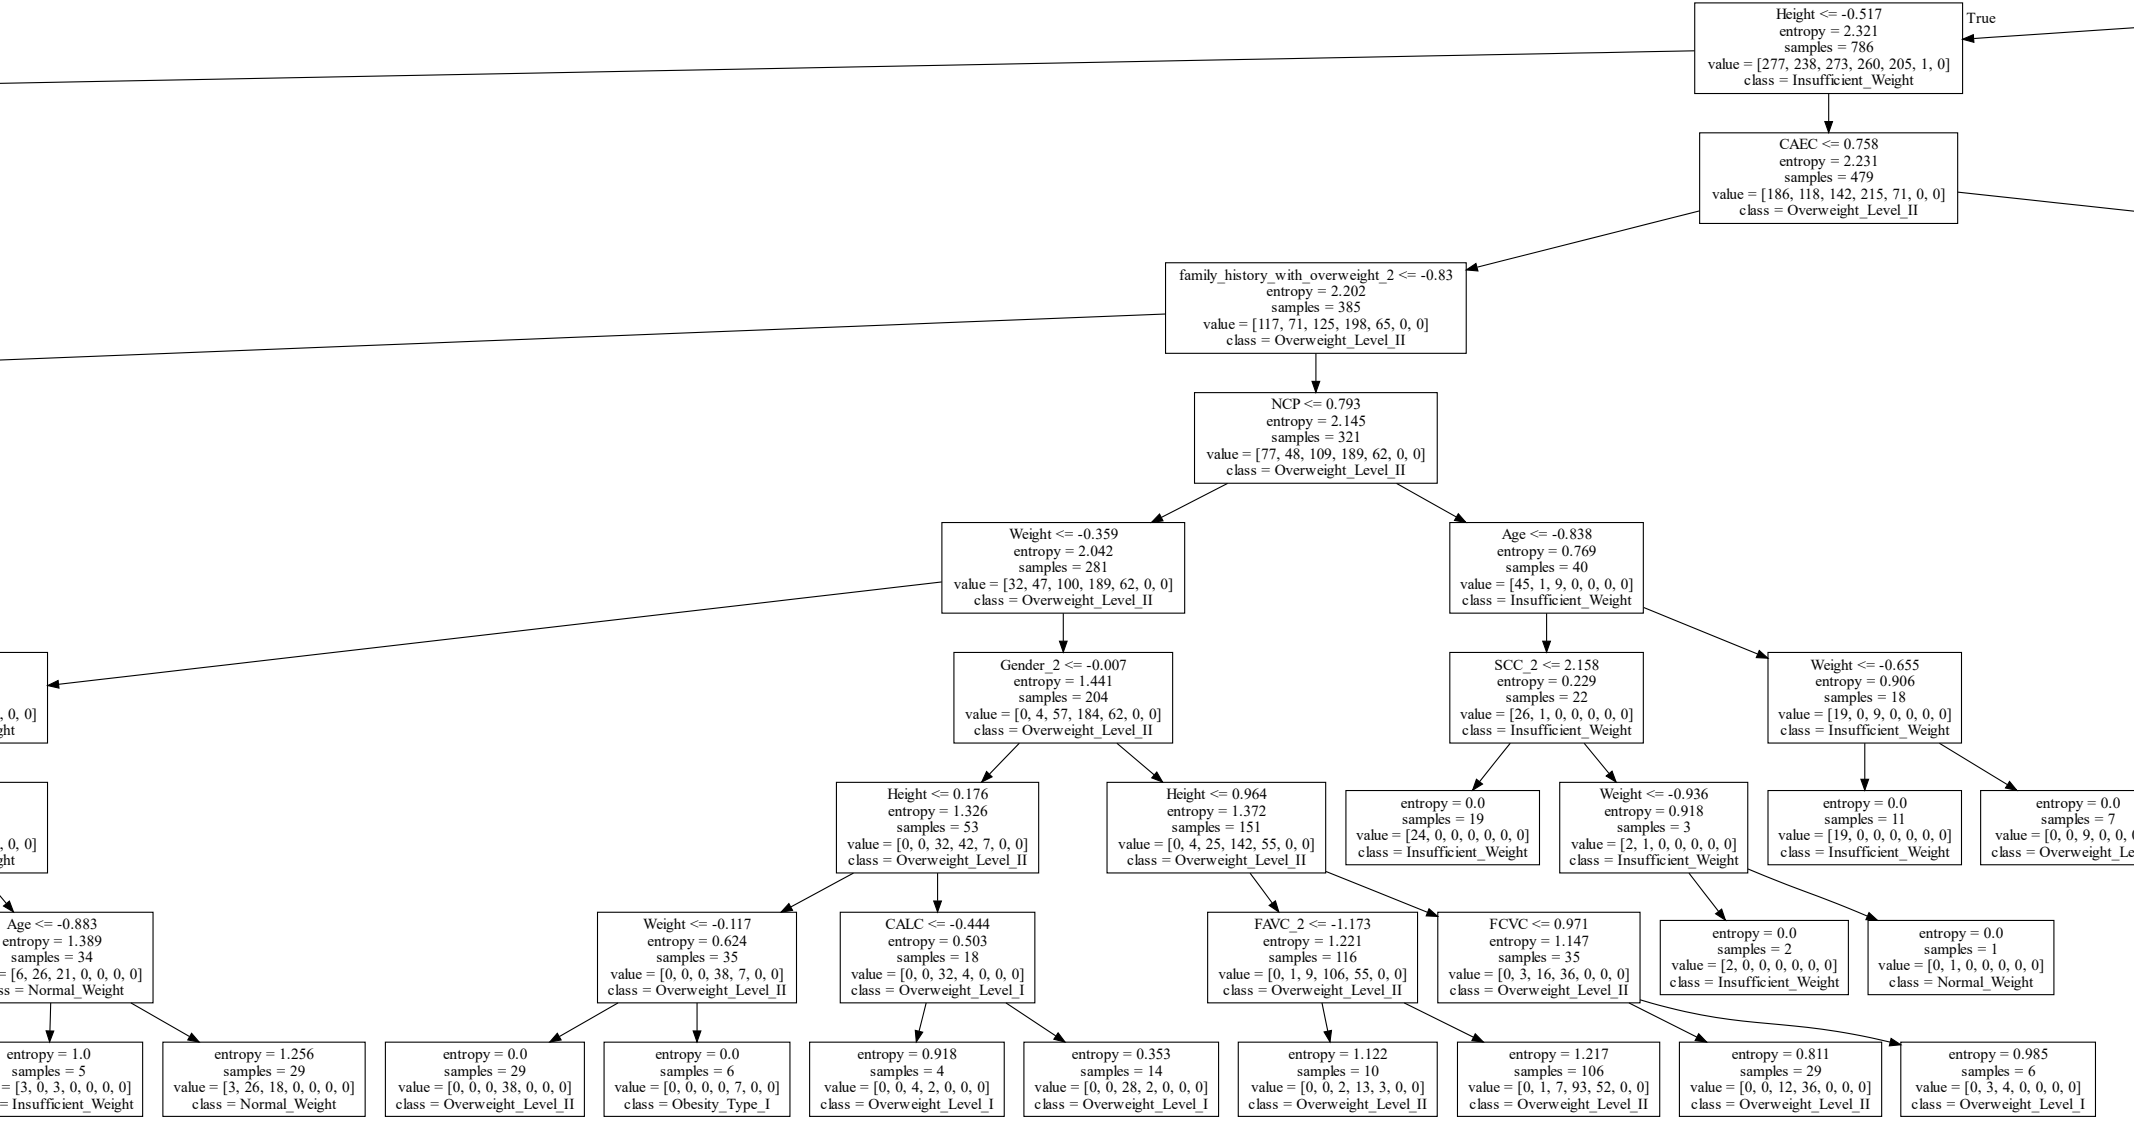
Height <= -0.517
entropy = 2.321
samples = 786
value = [277, 238, 273, 260, 205, 1, 0]
class = Insufficient_Weight
True
CAEC <= 0.758
entropy = 2.231
samples = 479
value = [186, 118, 142, 215, 71, 0, 0]
class = Overweight_Level_II
family_history_with_overweight_2 <= -0.83
entropy = 2.202
samples = 385
value = [117, 71, 125, 198, 65, 0, 0]
class = Overweight_Level_II
NCP <= 0.793
entropy = 2.145
samples = 321
value = [77, 48, 109, 189, 62, 0, 0]
class = Overweight_Level_II
Weight <= -0.359
entropy = 2.042
samples = 281
value = [32, 47, 100, 189, 62, 0, 0]
class = Overweight_Level_II
Age <= -0.838
entropy = 0.769
samples = 40
value = [45, 1, 9, 0, 0, 0, 0]
class = Insufficient_Weight
Gender_2 <= -0.007
entropy = 1.441
samples = 204
value = [0, 4, 57, 184, 62, 0, 0]
class = Overweight_Level_II
SCC_2 <= 2.158
entropy = 0.229
samples = 22
value = [26, 1, 0, 0, 0, 0, 0]
class = Insufficient_Weight
Weight <= -0.655
entropy = 0.906
samples = 18
value = [19, 0, 9, 0, 0, 0, 0]
class = Insufficient_Weight
Height <= 0.176
entropy = 1.326
samples = 53
value = [0, 0, 32, 42, 7, 0, 0]
class = Overweight_Level_II
Height <= 0.964
entropy = 1.372
samples = 151
value = [0, 4, 25, 142, 55, 0, 0]
class = Overweight_Level_II
entropy = 0.0
samples = 19
value = [24, 0, 0, 0, 0, 0, 0]
class = Insufficient_Weight
Weight <= -0.936
entropy = 0.918
samples = 3
value = [2, 1, 0, 0, 0, 0, 0]
class = Insufficient_Weight
entropy = 0.0
samples = 11
value = [19, 0, 0, 0, 0, 0, 0]
class = Insufficient_Weight
entropy = 0.0
samples = 7
value = [0, 0, 9, 0, 0,
class = Overweight_Le
Age <= -0.883
entropy = 1.389
samples = 34
= [6, 26, 21, 0, 0, 0, 0]
ss = Normal_Weight
Weight <= -0.117
entropy = 0.624
samples = 35
value = [0, 0, 0, 38, 7, 0, 0]
class = Overweight_Level_II
CALC <= -0.444
entropy = 0.503
samples = 18
value = [0, 0, 32, 4, 0, 0, 0]
class = Overweight_Level_I
FAVC_2 <= -1.173
entropy = 1.221
samples = 116
value = [0, 1, 9, 106, 55, 0, 0]
class = Overweight_Level_II
FCVC <= 0.971
entropy = 1.147
samples = 35
value = [0, 3, 16, 36, 0, 0, 0]
class = Overweight_Level_II
entropy = 0.0
samples = 2
value = [2, 0, 0, 0, 0, 0, 0]
class = Insufficient_Weight
entropy = 0.0
samples = 1
value = [0, 1, 0, 0, 0, 0, 0]
class = Normal_Weight
entropy = 1.0
samples = 5
= [3, 0, 3, 0, 0, 0, 0]
= Insufficient_Weight
entropy = 1.256
samples = 29
value = [3, 26, 18, 0, 0, 0, 0]
class = Normal_Weight
entropy = 0.0
samples = 29
value = [0, 0, 0, 38, 0, 0, 0]
class = Overweight_Level_II
entropy = 0.0
samples = 6
value = [0, 0, 0, 0, 7, 0, 0]
class = Obesity_Type_I
entropy = 0.918
samples = 4
value = [0, 0, 4, 2, 0, 0, 0]
class = Overweight_Level_I
entropy = 0.353
samples = 14
value = [0, 0, 28, 2, 0, 0, 0]
class = Overweight_Level_I
entropy = 1.122
samples = 10
value = [0, 0, 2, 13, 3, 0, 0]
class = Overweight_Level_II
entropy = 1.217
samples = 106
value = [0, 1, 7, 93, 52, 0, 0]
class = Overweight_Level_II
entropy = 0.811
samples = 29
value = [0, 0, 12, 36, 0, 0, 0]
class = Overweight_Level_II
entropy = 0.985
samples = 6
value = [0, 3, 4, 0, 0, 0, 0]
class = Overweight_Level_I

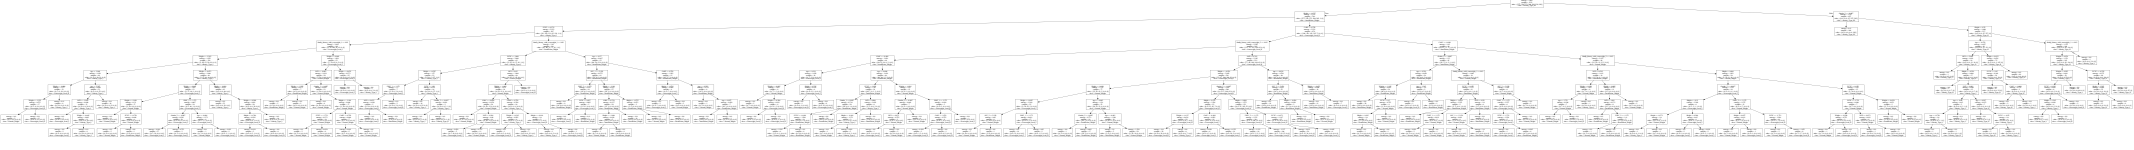

Tutorial

Tutorial Notebook

Reference Material

1. Chapter 6, 7 Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow by Aurélien Géron
2. https://ieeexplore.ieee.org/abstract/document/598994 – Random Forest academic paper.
3. https://towardsdatascience.com/understanding-random-forest-58381e0602d2
4. https://github.com/kjw0612/awesome-random-forest - Projects that have used Random Forests to achieve amazing solutions.
5. https://www.quora.com/What-are-the-advantages-and-disadvantages-for-a-random-forest-algorithm - Discussions on the advantages and disadvantages of Random Forests
6. https://www.youtube.com/watch?v=7VeUPuFGJHk - StatQuest: Decision Trees
7. https://www.youtube.com/watch?v=J4Wdy0Wc_xQ - StatQuest: Random Forests Part 1 - Building, Using and Evaluating

Thank you!

Questions?

The future of farmer financing

www.emata.ug